

# Introduction to Keras and TensorFlow

## All-ones or all-zeros tensors

```
In [ ]: import tensorflow as tf
x = tf.ones(shape=(2, 1))
print(x)

tf.Tensor(
[[1.]
 [1.]], shape=(2, 1), dtype=float32)
```

```
In [ ]: x = tf.zeros(shape=(2, 1))
print(x)

tf.Tensor(
[[0.]
 [0.]], shape=(2, 1), dtype=float32)
```

## Random tensors

```
In [ ]: x = tf.random.normal(shape=(3, 1), mean=0., stddev=1.)
print(x)

tf.Tensor(
[[0.8309784 ]
 [0.02322833]
 [0.20062953]], shape=(3, 1), dtype=float32)
```

```
In [ ]: x = tf.random.uniform(shape=(3, 1), minval=0., maxval=1.)
print(x)

tf.Tensor(
[[0.68336225]
 [0.27214646]
 [0.9351466 ]], shape=(3, 1), dtype=float32)
```

## NumPy arrays are assignable

```
In [ ]: import numpy as np
x = np.ones(shape=(2, 2))
x[0, 0] = 0.
```

## Creating a TensorFlow variable

```
In [ ]: v = tf.Variable(initial_value=tf.random.normal(shape=(3, 1)))
```

```
print(v)
```

```
<tf.Variable 'Variable:0' shape=(3, 1) dtype=float32, numpy=
array([[ -0.35847458],
       [  0.3193509 ],
       [  0.41017598]], dtype=float32)>
```

## Assigning a value to a TensorFlow variable

```
In [ ]: v.assign(tf.ones((3, 1)))
```

```
Out[ ]: <tf.Variable 'UnreadVariable' shape=(3, 1) dtype=float32, nu
numpy=
array([[1.],
       [1.],
       [1.]], dtype=float32)>
```

## Assigning a value to a subset of a TensorFlow variable

```
In [ ]: v[0, 0].assign(3.)
```

```
Out[ ]: <tf.Variable 'UnreadVariable' shape=(3, 1) dtype=float32, nu
numpy=
array([[3.],
       [1.],
       [1.]], dtype=float32)>
```

## Using assign\_add

```
In [ ]: v.assign_add(tf.ones((3, 1)))
```

```
Out[ ]: <tf.Variable 'UnreadVariable' shape=(3, 1) dtype=float32, nu
numpy=
array([[4.],
       [2.],
       [2.]], dtype=float32)>
```

## Tensor operations: Doing math in TensorFlow

### A few basic math operations

```
In [ ]: a = tf.ones((2, 2))
        b = tf.square(a)
        c = tf.sqrt(a)
        d = b + c
        e = tf.matmul(a, b)
        e *= d
```

## A second look at the GradientTape API

### Using the GradientTape

```
In [ ]: input_var = tf.Variable(initial_value=3.)
        with tf.GradientTape() as tape:
            result = tf.square(input_var)
            gradient = tape.gradient(result, input_var)
```

### Using GradientTape with constant tensor inputs

```
In [ ]: input_const = tf.constant(3.)
        with tf.GradientTape() as tape:
            tape.watch(input_const)
            result = tf.square(input_const)
            gradient = tape.gradient(result, input_const)
```

### Using nested gradient tapes to compute second-order gradients

```
In [ ]: time = tf.Variable(0.)
        with tf.GradientTape() as outer_tape:
            with tf.GradientTape() as inner_tape:
                position = 4.9 * time ** 2
                speed = inner_tape.gradient(position, time)
                acceleration = outer_tape.gradient(speed, time)
```

## An end-to-end example: A linear classifier in pure TensorFlow

### Generating two classes of random points in a 2D plane

```
In [ ]: num_samples_per_class = 1000
        negative_samples = np.random.multivariate_normal(
            mean=[0, 3],
            cov=[[1, 0.5],[0.5, 1]],
            size=num_samples_per_class)
        positive_samples = np.random.multivariate_normal(
            mean=[3, 0],
            cov=[[1, 0.5],[0.5, 1]],
            size=num_samples_per_class)
```

### Stacking the two classes into an array with shape (2000, 2)


```
In [ ]: inputs = np.vstack((negative_samples, positive_samples)).astype(float)
```

### Generating the corresponding targets (0 and 1)

```
In [ ]: targets = np.vstack((np.zeros((num_samples_per_class, 1), dtype=int),
                                np.ones((num_samples_per_class, 1), dtype=int)))
```

## Plotting the two point classes

```
In [ ]: import matplotlib.pyplot as plt
plt.scatter(inputs[:, 0], inputs[:, 1], c=targets[:, 0])
plt.show()
```



## Creating the linear classifier variables

```
In [ ]: input_dim = 2
output_dim = 1
W = tf.Variable(initial_value=tf.random.uniform(shape=(input_dim, output_dim)))
b = tf.Variable(initial_value=tf.zeros(shape=(output_dim,)))
```

## The forward pass function

```
In [ ]: def model(inputs):
    return tf.matmul(inputs, W) + b
```

## The mean squared error loss function

```
In [ ]: def square_loss(targets, predictions):
    per_sample_losses = tf.square(targets - predictions)
    return tf.reduce_mean(per_sample_losses)
```

## The training step function

```
In [ ]: learning_rate = 0.1

def training_step(inputs, targets):
    with tf.GradientTape() as tape:
        predictions = model(inputs)
```

```
    loss = square_loss(targets, predictions)
    grad_loss_wrt_W, grad_loss_wrt_b = tape.gradient(loss, [
W.assign_sub(grad_loss_wrt_W * learning_rate)
b.assign_sub(grad_loss_wrt_b * learning_rate)
return loss
```

## The batch training loop


```
In [ ]: for step in range(40):
        loss = training_step(inputs, targets)
        print(f"Loss at step {step}: {loss:.4f}")
```

```
Loss at step 0: 1.8413
Loss at step 1: 0.3466
Loss at step 2: 0.1495
Loss at step 3: 0.1168
Loss at step 4: 0.1059
Loss at step 5: 0.0983
Loss at step 6: 0.0917
Loss at step 7: 0.0858
Loss at step 8: 0.0804
Loss at step 9: 0.0754
Loss at step 10: 0.0709
Loss at step 11: 0.0668
Loss at step 12: 0.0631
Loss at step 13: 0.0597
Loss at step 14: 0.0566
Loss at step 15: 0.0538
Loss at step 16: 0.0512
Loss at step 17: 0.0489
Loss at step 18: 0.0467
Loss at step 19: 0.0448
Loss at step 20: 0.0430
Loss at step 21: 0.0414
Loss at step 22: 0.0399
Loss at step 23: 0.0386
Loss at step 24: 0.0373
Loss at step 25: 0.0362
Loss at step 26: 0.0352
Loss at step 27: 0.0343
Loss at step 28: 0.0334
Loss at step 29: 0.0327
Loss at step 30: 0.0320
Loss at step 31: 0.0313
Loss at step 32: 0.0307
Loss at step 33: 0.0302
Loss at step 34: 0.0297
Loss at step 35: 0.0293
Loss at step 36: 0.0289
Loss at step 37: 0.0285
```

Loss at step 38: 0.0282


Loss at step 39: 0.0279

```
In [ ]: predictions = model(inputs)
plt.scatter(inputs[:, 0], inputs[:, 1], c=predictions[:, 0])
plt.show()
```



```
In [ ]: x = np.linspace(-1, 4, 100)
y = - W[0] / W[1] * x + (0.5 - b) / W[1]
plt.plot(x, y, "-r")
plt.scatter(inputs[:, 0], inputs[:, 1], c=predictions[:, 0])
```

Out[ ]: <matplotlib.collections.PathCollection at 0x7fa2db6dad10>



**Anatomy of a neural network:**

# Understanding core Keras APIs

## Layers: The building blocks of deep learning

### The base Layer class in Keras

#### A Dense layer implemented as a Layer subclass

```
In [ ]: from tensorflow import keras

class SimpleDense(keras.layers.Layer):

    def __init__(self, units, activation=None):
        super().__init__()
        self.units = units
        self.activation = activation

    def build(self, input_shape):
        input_dim = input_shape[-1]
        self.W = self.add_weight(shape=(input_dim, self.units),
                                initializer="random_normal")
        self.b = self.add_weight(shape=(self.units,),
                                initializer="zeros")

    def call(self, inputs):
        y = tf.matmul(inputs, self.W) + self.b
        if self.activation is not None:
            y = self.activation(y)
        return y
```

```
In [ ]: my_dense = SimpleDense(units=32, activation=tf.nn.relu)
input_tensor = tf.ones(shape=(2, 784))
output_tensor = my_dense(input_tensor)
print(output_tensor.shape)

(2, 32)
```

### Automatic shape inference: Building layers on the fly

```
In [ ]: from tensorflow.keras import layers
layer = layers.Dense(32, activation="relu")
```

```
In [ ]: from tensorflow.keras import models
from tensorflow.keras import layers
model = models.Sequential([
    layers.Dense(32, activation="relu"),
```

```
layers.Dense(32)
])
```

```
In [ ]: model = keras.Sequential([
        SimpleDense(32, activation="relu"),
        SimpleDense(64, activation="relu"),
        SimpleDense(32, activation="relu"),
        SimpleDense(10, activation="softmax")
    ])
```

## From layers to models

### The "compile" step: Configuring the learning process

```
In [ ]: model = keras.Sequential([keras.layers.Dense(1)])
model.compile(optimizer="rmsprop",
              loss="mean_squared_error",
              metrics=["accuracy"])
```

```
In [ ]: model.compile(optimizer=keras.optimizers.RMSprop(),
                      loss=keras.losses.MeanSquaredError(),
                      metrics=[keras.metrics.BinaryAccuracy()])
```

## Picking a loss function

## Understanding the fit() method

### Calling fit() with NumPy data

```
In [ ]: history = model.fit(
        inputs,
        targets,
        epochs=5,
        batch_size=128
    )
```

Epoch 1/5

16/16 [=====] - 1s 2ms/step - loss: 15.0589 - binary\_accuracy: 0.0020

Epoch 2/5

16/16 [=====] - 0s 2ms/step - loss: 14.6048 - binary\_accuracy: 0.0020

Epoch 3/5

16/16 [=====] - 0s 2ms/step - loss: 14.2170 - binary\_accuracy: 0.0020



```
Epoch 4/5
16/16 [=====] - 0s 2ms/step - loss:
13.8422 - binary_accuracy: 0.0020
Epoch 5/5
16/16 [=====] - 0s 2ms/step - loss:
13.4738 - binary_accuracy: 0.0020
```

```
In [ ]: history.history
```

```
Out[ ]: {'loss': [15.058869361877441,
 14.604792594909668,
 14.217011451721191,
 13.84217357635498,
 13.473837852478027],
'binary_accuracy': [0.002000000949949026,
 0.002000000949949026,
 0.002000000949949026,
 0.002000000949949026,
 0.002000000949949026]}
```

## Monitoring loss and metrics on validation data

### Using the `validation_data` argument

```
In [ ]: model = keras.Sequential([keras.layers.Dense(1)])
model.compile(optimizer=keras.optimizers.RMSprop(learning_rate=0.01),
              loss=keras.losses.MeanSquaredError(),
              metrics=[keras.metrics.BinaryAccuracy()])

indices_permutation = np.random.permutation(len(inputs))
shuffled_inputs = inputs[indices_permutation]
shuffled_targets = targets[indices_permutation]

num_validation_samples = int(0.3 * len(inputs))
val_inputs = shuffled_inputs[:num_validation_samples]
val_targets = shuffled_targets[:num_validation_samples]
training_inputs = shuffled_inputs[num_validation_samples:]
training_targets = shuffled_targets[num_validation_samples:]

model.fit(
    training_inputs,
    training_targets,
    epochs=5,
    batch_size=16,
    validation_data=(val_inputs, val_targets)
)
```

```
Epoch 1/5
88/88 [=====] - 1s 5ms/step - loss:
```

```

0.3618 - binary_accuracy: 0.8979 - val_loss: 0.0778 - val_bi
nary_accuracy: 0.9450
Epoch 2/5
88/88 [=====] - 0s 3ms/step - loss:
0.0765 - binary_accuracy: 0.9536 - val_loss: 0.0733 - val_bi
nary_accuracy: 0.9617
Epoch 3/5
88/88 [=====] - 0s 3ms/step - loss:
0.0721 - binary_accuracy: 0.9600 - val_loss: 0.0659 - val_bi
nary_accuracy: 0.9650
Epoch 4/5
88/88 [=====] - 0s 3ms/step - loss:
0.0687 - binary_accuracy: 0.9564 - val_loss: 0.0287 - val_bi
nary_accuracy: 0.9967
Epoch 5/5
88/88 [=====] - 0s 3ms/step - loss:
0.0708 - binary_accuracy: 0.9657 - val_loss: 0.0350 - val_bi
nary_accuracy: 1.0000

```

```
Out[ ]: <keras.callbacks.History at 0x7fa2d9b5ce90>
```

## Inference: Using a model after training

```
In [ ]: predictions = model.predict(val_inputs, batch_size=128)
print(predictions[:10])
```

```

5/5 [=====] - 0s 3ms/step
[[0.24031283]
 [0.84266007]
 [0.91847444]
 [0.14388952]
 [0.11975732]
 [0.08451074]
 [0.17176777]
 [0.09675378]
 [0.14447355]
 [0.07130358]]

```