

```
In [ ]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python docker image: https://github.com/kaggle/docker-
# For example, here's several helpful packages to load in

import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list all

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# Any results you write to the current directory are saved as output.
```

```
In [1]: import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout, Conv2D, MaxPooling2D, B
from tensorflow.keras.utils import to_categorical, plot_model

import matplotlib.pyplot as plt
```

```
In [2]: train = pd.read_csv('/kaggle/input/digit-recognizer/train.csv')
test = pd.read_csv('/kaggle/input/digit-recognizer/test.csv')
```

```
In [3]: train.head(2)
```

```
Out[3]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	...	pixel774	pixel77
0	1	0	0	0	0	0	0	0	0	0	...	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	

2 rows × 785 columns



```
In [4]: test.head(2)
```

```
Out[4]:
```

	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel774	pixel77
0	0	0	0	0	0	0	0	0	0	0	...	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	

2 rows × 784 columns



```
In [7]: train = np.array(train)
test = np.array(test)

train_x = train[:, 1:]
train_y = train[:, 0]

# Normalize the data
```

```
train_x = train_x / 255.0
test = test / 255.0
```

In [8]: train_x

```
Out[8]: array([[0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]])
```

In [10]: *#Reshape*
train_x = train_x.reshape(-1,28,28,1)
test = test.reshape(-1,28,28,1)

In [12]: train_y

```
Out[12]: array([1, 0, 1, ..., 7, 6, 9])
```

In [13]: train_y = to_categorical(train_y)


In [14]: train_y

```
Out[14]: array([[0., 1., 0., ..., 0., 0., 0.],
               [1., 0., 0., ..., 0., 0., 0.],
               [0., 1., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 1., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 1.]], dtype=float32)
```

In [15]: print(train_x.shape)

```
(42000, 28, 28, 1)
```

In [19]: g = plt.imshow(train_x[1][:,:,0])



In [22]: model = Sequential()

model.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',
 activation = 'relu', input_shape = (28,28,1)))
model.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',
 activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

```

model.add(Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',
                 activation = 'relu'))
model.add(Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',
                 activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256, activation = "relu"))
model.add(Dropout(0.5))
model.add(Dense(10, activation = "softmax"))

```

```

In [27]: from tensorflow.keras.optimizers import RMSprop
        from tensorflow.keras.preprocessing.image import ImageDataGenerator
        from tensorflow.keras.callbacks import ReduceLROnPlateau
        #Trying RMSProp optimizer

        optimizer_rms = RMSprop(lr=0.001, rho=0.9, epsilon=1e-08, decay=0.0)
        # Compile the model
        model.compile(optimizer = optimizer_rms , loss = "categorical_crossentropy", metrics

```

```

In [28]: BATCH_SIZE=64
        EPOCH=5

```

```

In [29]: history = model.fit(train_x, train_y, batch_size=BATCH_SIZE, epochs=EPOCH, validation

Train on 33600 samples, validate on 8400 samples
Epoch 1/5
33600/33600 [=====] - 126s 4ms/sample - loss: 0.2354 - accu
racy: 0.9255 - val_loss: 0.0603 - val_accuracy: 0.9800
Epoch 2/5
33600/33600 [=====] - 124s 4ms/sample - loss: 0.0771 - accu
racy: 0.9762 - val_loss: 0.0441 - val_accuracy: 0.9873
Epoch 3/5
33600/33600 [=====] - 124s 4ms/sample - loss: 0.0578 - accu
racy: 0.9822 - val_loss: 0.0423 - val_accuracy: 0.9873
Epoch 4/5
33600/33600 [=====] - 124s 4ms/sample - loss: 0.0507 - accu
racy: 0.9854 - val_loss: 0.0366 - val_accuracy: 0.9894
Epoch 5/5
33600/33600 [=====] - 124s 4ms/sample - loss: 0.0450 - accu
racy: 0.9870 - val_loss: 0.0354 - val_accuracy: 0.9907

```

```

In [31]: from tensorflow.keras.optimizers import Adam
        optimizer_adam = Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07,n
        # Compile the model
        model.compile(optimizer = optimizer_adam , loss = "categorical_crossentropy", metric
        # Since target variable has multiple classes, loss fnctn can be categorical cross en
        history_adam = model.fit(train_x, train_y, batch_size=BATCH_SIZE, epochs=EPOCH, vali

Train on 33600 samples, validate on 8400 samples
Epoch 1/5
33600/33600 [=====] - 123s 4ms/sample - loss: 0.0391 - accu
racy: 0.9879 - val_loss: 0.0333 - val_accuracy: 0.9895
Epoch 2/5
33600/33600 [=====] - 123s 4ms/sample - loss: 0.0352 - accu
racy: 0.9895 - val_loss: 0.0350 - val_accuracy: 0.9902
Epoch 3/5
33600/33600 [=====] - 124s 4ms/sample - loss: 0.0324 - accu
racy: 0.9900 - val_loss: 0.0328 - val_accuracy: 0.9913
Epoch 4/5
33600/33600 [=====] - 123s 4ms/sample - loss: 0.0278 - accu
racy: 0.9917 - val_loss: 0.0348 - val_accuracy: 0.9910

```

Epoch 5/5
33600/33600 [=====] - 129s 4ms/sample - loss: 0.0260 - accuracy: 0.9921 - val_loss: 0.0275 - val_accuracy: 0.9921

```
In [32]: #Stochastic gradient with momentum and Nesterov  
#we can slo use momentum optimzer (sgd+momentum)  
from tensorflow.keras.optimizers import SGD  
optimizer_sgd=SGD(learning_rate=0.01, momentum=0.9, nesterov=True, name='SGD')  
model.compile(optimizer = optimizer_sgd , loss = "categorical_crossentropy", metrics  
history_sgd = model.fit(train_x, train_y, batch_size=BATCH_SIZE, epochs=EPOCH, valid
```

Train on 33600 samples, validate on 8400 samples

Epoch 1/5
33600/33600 [=====] - 124s 4ms/sample - loss: 0.0180 - accuracy: 0.9943 - val_loss: 0.0333 - val_accuracy: 0.9927
Epoch 2/5
33600/33600 [=====] - 123s 4ms/sample - loss: 0.0156 - accuracy: 0.9953 - val_loss: 0.0410 - val_accuracy: 0.9927
Epoch 3/5
33600/33600 [=====] - 123s 4ms/sample - loss: 0.0113 - accuracy: 0.9965 - val_loss: 0.0361 - val_accuracy: 0.9931
Epoch 4/5
33600/33600 [=====] - 134s 4ms/sample - loss: 0.0124 - accuracy: 0.9958 - val_loss: 0.0382 - val_accuracy: 0.9931
Epoch 5/5
33600/33600 [=====] - 123s 4ms/sample - loss: 0.0099 - accuracy: 0.9965 - val_loss: 0.0328 - val_accuracy: 0.9931

SGD Optimizer proves to be the best one, Loss consistently reduced during each iteration